

# Assignment N° 4: Transformer with PyTorch

Daniel Jader Pellattiero

고려대학교

Seoul, Republic of Korea

danieljaderpellattiero@outlook.com

## Abstract

*The goal of this assignment is to develop a Transformer to perform classification on a dataset of IMDb reviews, consisting of both 15k training samples and 10k for validation and testing.*

## 1. Implementation

### 1.1. Positional encoding

To supply the model with sequential-order information that pure token embeddings lack, we add a deterministic sinusoidal signal to each representation.

For position  $p \in \{0, \dots, L - 1\}$  and embedding coordinate  $k \in \{0, \dots, d_{\text{model}} - 1\}$ ,

$$\text{PE}_{p,k} = \begin{cases} \sin(p \cdot 10000^{-k/d_{\text{model}}}), & k \text{ even}, \\ \cos(p \cdot 10000^{-(k-1)/d_{\text{model}}}), & k \text{ odd}. \end{cases}$$

These sinusoids give every position a unique phase vector while ensuring that any fixed shift in  $p$  can be expressed as a linear function of the encodings. In the forward pass we add the  $\text{PE}_{p,k}$  to the learned token embeddings, followed by dropout, leaving dimensionality unchanged.

### 1.2. Multi Head Attention

Scaled dot-product attention projects the query  $Q$ , key  $K$  and value  $V$  tensors into  $h$  parallel subspaces of dimension  $d_k = d_{\text{model}}/h$ . Each head computes:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V,$$

optionally masking illegal positions before softmax. The  $h$  outputs are concatenated and mapped back to the model dimension by a learned matrix  $W_O$ . Splitting attention this way lets the network attend to multiple relational patterns simultaneously without increasing asymptotic complexity beyond  $\mathcal{O}(L^2 d_{\text{model}})$ .

### 1.3. Encoder

A single encoder block alternates self-attention and position-wise feed-forward sub-layers, each wrapped in residual connections and layer normalisation. Denoting the input by  $x$ , we have:

$$\begin{aligned} x' &= x + \text{Drop}(\text{MultiHead}(x)), & \tilde{x} &= \text{LayerNorm}(x') \\ y' &= \tilde{x} + \text{Drop}(W_2 \sigma(W_1 \tilde{x})), & y &= \text{LayerNorm}(y') \end{aligned}$$

where  $W_1$  and  $W_2$  form a two-layer MLP of hidden size  $d_{\text{ff}}$  and  $\sigma$  is either ReLU or GELU. This architecture enables global context mixing followed by non-linear feature transformation, while residual paths stabilise gradients in deep stacks.

The full encoder is a stack of  $N$  identical layers that share the structure above but not parameters:

$$x^{(0)} = \text{Embed}(w) + \text{PE},$$

$$x^{(l+1)} = \text{EncoderLayer}^{(l)}(x^{(l)}), \quad l = 0, \dots, N - 1.$$

An optional final layer normalisation yields  $x^{(N)}$ , a contextualised sequence representation suitable for classification or other tasks. Depth  $N$  and the feed-forward width  $d_{\text{ff}}$  control model capacity, while the number of heads  $h$  governs how many relational patterns can be modelled in parallel.

## 2. Experimental results

The implemented model, characterized by a 4-level encoder, 12 attention heads and a dimensionality of 768, showed an accuracy value of 62.38% on the test set after one epoch of training with the Adam optimizer (learning rate of  $5e - 5$ ) and Cross Entropy as a loss function.

Using the same training modes, the BERT model (pre-trained) demonstrated to have a higher accuracy of 82.66%.