

# 1<sup>st</sup> Assignment: MLP Implementation

Daniel Jader Pellattiero

Korea University

danieljaderpellattiero@gmail.com

## Abstract

*The goal of the first assignment is to implement a two-layer neural network that performs image classification on the MNIST data set.*

## 1. Introduction

The MNIST data set consists of  $28 \times 28$  grayscale images of 70K handwritten digits ranging from (0–9). The goal is to train a classifier that correctly recognizes which digit is shown in each image. I have implemented a basic two-layer Multi-Layer Perceptron (MLP) with one hidden layer of size  $H$ :

$$\mathbf{z}_1 = X\mathbf{W}_1 + \mathbf{b}_1, \quad \mathbf{a}_1 = \text{ReLU}(\mathbf{z}_1), \quad \mathbf{z}_2 = \mathbf{a}_1\mathbf{W}_2 + \mathbf{b}_2, \quad (1)$$

where  $X$  is the input matrix (flattened images),  $\mathbf{W}_1, \mathbf{W}_2$  are the weight matrices,  $\mathbf{b}_1, \mathbf{b}_2$  are the biases, and ReLU is the activation function of the hidden layer.

**Loss Function** I compute the various classes' probabilities via softmax:

$$\text{softmax}(z_i)_c = \frac{\exp(z_{i,c})}{\sum_{k=1}^C \exp(z_{i,k})}, \quad (2)$$

and then I use the cross-entropy loss:

$$L_{data} = -\frac{1}{N} \sum_{i=1}^N \log(p_{i,y_i}) \quad (3)$$

where  $p_{i,y_i} = \text{softmax}(z_i)_{y_i}$ . I also add an L2 penalty to our weight parameters  $\mathbf{W}_1, \mathbf{W}_2$  like so:

$$L_{reg} = -\frac{\lambda}{2} (\|\mathbf{W}_1\|_2^2 + \|\mathbf{W}_2\|_2^2) \quad (4)$$

**Train Function** During training, I perform mini-batch Stochastic Gradient Descent (SGD). I firstly sample a random batch of images and labels and then I compute the

forward pass to obtain the loss. Next, I perform back-propagation to calculate gradients w.r.t. all parameters— $\mathbf{W}_1, \mathbf{b}_1, \mathbf{W}_2, \mathbf{b}_2$ . Finally, I update these parameters via  $\theta \leftarrow \theta - \eta \nabla_{\theta} L$ , where  $\eta$  is the learning rate.

## 2. Methodology

**First configuration** Follows the initial configuration of the model's hyperparameters:

- Input size: (28; 28)
- Hidden size: 100
- Iterations: 1000
- Batch size: 200
- Learning rate:  $1e^{-3}$
- Learning rate decay: 0.95
- Regularization:  $1e^{-3}$

The performance of this first version of the neural network was not satisfactory as it reached a validation accuracy of 0.105 on the validation set. To address this, I automatically tested various values for the hidden layer size ( $H$ ), learning rate ( $\eta$ ), number of epochs, and regularization strength ( $\lambda$ ). At the end I achieved a classification accuracy of over 39% on the validation set and over 50% on the test set, which is better than the baseline.

## 3. Performance and Hyperparameter Tuning

The most promising configurations tested, through an automated script, are reported below.

Model	$H$	$\eta$	$\lambda$	Val. Acc.
(Best)	300	1e-1	1e-4	0.940
Model B	100	1e-1	1e-4	0.939
Model C	300	1e-1	1e-3	0.939
(etc.)				

The batch size and learning rate decay remained the same, while the number of iterations was doubled.